


```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: dt=pd.read_csv(r"C:\Users\91903\Downloads\fiat500_VehicleSelection_Dataset.csv")
dt
```

```
Out[2]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611560
1	2	pop	51	1186	32500	1	45.666359	12.241890
2	3	sport	74	4658	142228	1	45.503300	11.417840
3	4	lounge	51	2739	160000	1	40.633171	17.634609
4	5	pop	73	3074	106880	1	41.903221	12.495650
...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870
1535	1536	pop	51	2223	60457	1	45.481541	9.413480
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270
1537	1538	pop	51	1766	54276	1	40.323410	17.568270

1538 rows × 9 columns



```
In [5]: dt=dt[['engine_power','price']]
dt.columns=['Engine','Pric']
```


```
In [6]: dt.head(10)
```

```
Out[6]:
```

	Engine	Pric
0	51	8900
1	51	8800
2	74	4200
3	51	6000
4	73	5700
5	74	7900
6	51	10750
7	51	9190
8	73	5600
9	51	6000

```
In [8]: sns.lmplot(x='Engine',y='Pric',data=dt,order=2,ci=None)
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x1d7f5f418d0>
```



In [9]: `dt.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   Engine  1538 non-null    int64
 1   Pric    1538 non-null    int64
dtypes: int64(2)
memory usage: 24.2 KB
```

In [10]: `dt.describe()`

Out[10]:

	Engine	Pric
<b>count</b>	1538.000000	1538.000000
<b>mean</b>	51.904421	8576.003901
<b>std</b>	3.988023	1939.958641
<b>min</b>	51.000000	2500.000000
<b>25%</b>	51.000000	7122.500000
<b>50%</b>	51.000000	9000.000000
<b>75%</b>	51.000000	10000.000000
<b>max</b>	77.000000	11100.000000

In [11]: `dt.fillna(method='ffill')`

Out[11]:

	Engine	Pric
<b>0</b>	51	8900
<b>1</b>	51	8800
<b>2</b>	74	4200
<b>3</b>	51	6000
<b>4</b>	73	5700
...	...	...
<b>1533</b>	51	5200
<b>1534</b>	74	4600
<b>1535</b>	51	7500
<b>1536</b>	51	5990
<b>1537</b>	51	7900

1538 rows × 2 columns

```
In [12]: x=np.array(dt['Engine']).reshape(-1,1)
y=np.array(dt['Pric']).reshape(-1,1)
```

```
In [13]: dt.dropna(inplace=True)
```

C:\Users\91903\AppData\Local\Temp\ipykernel\_12044\735218168.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame


See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
dt.dropna(inplace=True)
```

```
In [14]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
reg=LinearRegression()
reg.fit(X_train,y_train)
print(reg.score(X_test,y_test))
```


```
0.09528364801504663
```

```
In [15]: y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color='b')
plt.plot(X_test,y_pred,color='k')
plt.show()
```




```
In [16]: dt500=dt[:][:500]  
sns.lmplot(x="Engine",y="Pric",data=dt500,order=1,ci=None)
```

Out[16]: <seaborn.axisgrid.FacetGrid at 0x1d7d1d47460>



```
In [17]: dt500.fillna(method='ffill',inplace=True)
X=np.array(dt500['Engine']).reshape(-1,1)
y=np.array(dt500['Pric']).reshape(-1,1)
dt500.dropna(inplace=True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
reg=LinearRegression()
reg.fit(X_train,y_train)
print("Regression:",reg.score(X_test,y_test))
y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color="b")
plt.plot(X_test,y_pred,color='k')
plt.show()
```

Regression: 0.08361221541270236



```
In [18]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.08361221541270236

#conclusion : Linear regression is not fit for the model

In [ ]: